# Reliable Data Transfer between Two Machines

Issues

- Synchronization
  - How does the receiver know when the sender is sending data?
- Encoding
  - Representing data with signals
- Error control
  - How to ensure the receiver gets the correct data
- Flow control
  - How to ensure the sender does not swamp the receiver (fast sender, slow receiver)

Synchronization

Need for synchronization

- Basic steps in transmitting digital data
  - Transmitter sends a stream of 0 or 1 bits
  - Receiver samples incoming signal once per bit time to see if it is a 0 or 1
    - Typically at the center of the bit
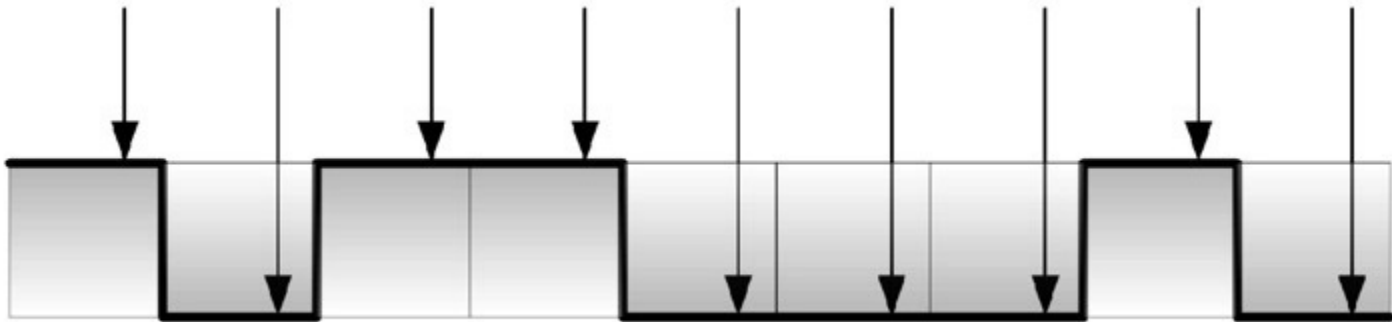
- Transmitter sends one bit (0 or 1) every 1 millisecond (say)
  - Say there is a clock that ticks every 1 millisecond
  - Transmitter puts a 1 or 0 on the line at each tick of the clock
    - bit 1 on tick 1, bit 2 on tick 2, …
- What should the receiver do to get the pattern correctly?
  - Need to know when to start counting (when will bit 1 arrive)
    - A-priori agreed or sender has to tell (how?)
  - Need to know when to look for the next successive bits (duration of one bit)

Framing

- Break up the bit pattern into multiple frames
- Issues
  - Needs identifier for each frame to distinguish between frames
  - Needs technique to identify where a frame ends and the next one begins
    - Frame synchronization problem

Frame Synchronization

- When data is transferred from the transmitter to the receiver, unless steps are taken to provide synchronization, the receiver may start interpreting the data erroneously

- Two common approaches:
  - Asynchronous Transmission
  - Synchronous Transmission

Asynchronous Transmission

- Transmitter and Receiver has separate clocks
- Data rate and frame format negotiated a-priori
- Data are transmitted one character at a time (5-8 bits)
  - Timing or synchronization must only be maintained within each character
  - Start bit at the beginning of each new character to resynchronize (start the receiver clock at the right time)
- There may be gaps between two successive character
- When no character is being transmitted, the line between transmitter and receiver is in an *idle* state
  - Start, stop, and parity bits.

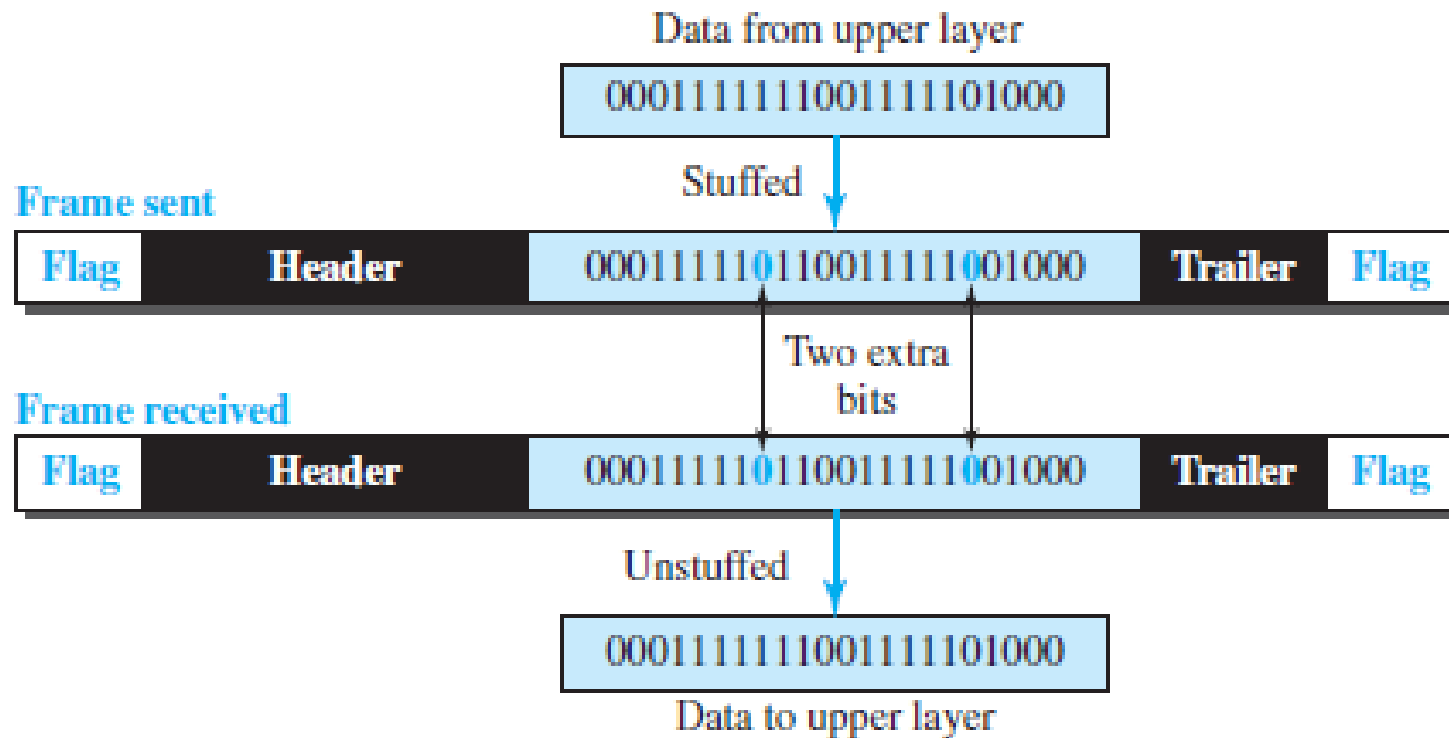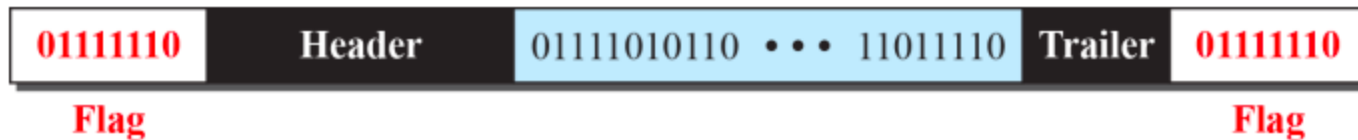| Start | | | | Data | | | | P | Stop | |

Synchronous Transmission

- Receiver clock is synchronized with transmitter clock

- A block of bits is transmitted in a steady stream without start and stop codes
  - The block may be arbitrarily long
  - To prevent timing drift between transmitter and receiver, clock signal is embedded in the data signal or separate clock signal sent

- There is another level of synchronization required, so as to allow the receiver to determine the beginning and end of a block of data
  - Every block begins with a preamble bit pattern, and generally ends with a postamble bit pattern

Synchronous Transmission (contd.)

- A typical synchronous frame format:
  - 8-bit flag (preamble)
  - Control fields
  - Data field
  - Control fields
  - 8-bit flag (postamble)
- Special inter-frame fill pattern is sent when no frame to send to maintain synchronization
- What if data itself contains the preamble/postamble pattern?
  - Use bit-stuffing

01111110	Header	01111010110 • • • 11011110	Trailer	01111110
Flag				Flag

Data from upper layer

00011111111001111101000

Stuffed

Frame sent

| Flag | Header | 000111110110011111001000 | Trailer | Flag |

Two extra bits

Frame received

| Flag | Header | 000111110110011111001000 | Trailer | Flag |

Unstuffed

00011111111001111101000

Data to upper layer

- For sizable blocks of data, synchronous transmission is far more efficient than asynchronous mode
  - Asynchronous transmission requires 20% or more overhead
  - The control information, preamble and postamble in synchronous transmission are typically less than 100 bits

Encoding Techniques

Data Encoding Techniques

- Data is transmitted by propagation and processing of signals

- Encoding – specifying how data is represented by signals
  - Digital data over digital signal
  - Digital data over analog signal

Encoding Digital Data with Digital Signals

- Digital signal
    - Uses discrete, discontinuous, voltage pulses
    - Each pulse is a signal element
    - Binary data is encoded into signal elements
- Data
    - Bit string of 0's and 1's
    - Each bit is present for a duration T (bit interval)
    - Data rate = 1 / T bps

Some Terminology

- Unipolar
  - All signal elements have same sign
- Polar
  - Two logic states represented by +ve/-ve voltage
- Duration or length of a bit
  - Time taken for transmitter to emit the bit
- Modulation rate
  - Rate at which the signal level changes
- Mark and Space (historical)
  - Binary 1 and Binary 0 respectively

Issues in Encoding

- Signal Spectrum
  - Lack of high frequencies reduces required bandwidth
  - Lack of dc component allows ac coupling via transformer, providing isolation reducing intereference
- Clocking issues
  - Synchronizing transmitter and receiver is essential
  - External clock is one way used for synchronization
  - Synchronizing mechanism based on signal is also used & preferred (over using an external clock)

- Error detection
  - Can be built into signal encoding
- Signal interference and noise immunity
  - Some codes are better than others
- Cost and complexity
  - Higher signal rate (& thus data rate) lead to higher costs
  - Some codes require signal rate greater than data rate

Some Encoding Schemes

- Nonreturn to Zero-Level (NRZ-L)

- Nonreturn to Zero Inverted (NRZI)

- Bipolar-AMI (Alternate Mark Inversion)

- Pseudoternary

- Manchester

- Differential Manchester

Nonreturn to Zero-Level (NRZ-L)

- Two different voltages for 0 and 1 bits
- Voltage held constant during bit interval
- Example:
  - 0 = High
  - 1 = Low

Nonreturn to Zero Inverted (NRZ-I)

- *Nonreturn to zero inverted on ones*
- Voltage held constant during bit interval
- Data encoded as presence or absence of signal transition at beginning of bit time
- Transition (low to high or high to low) denotes a binary 1
- No transition denotes binary 0
- An example of differential encoding (Data represented by changes rather than levels)

NRZ
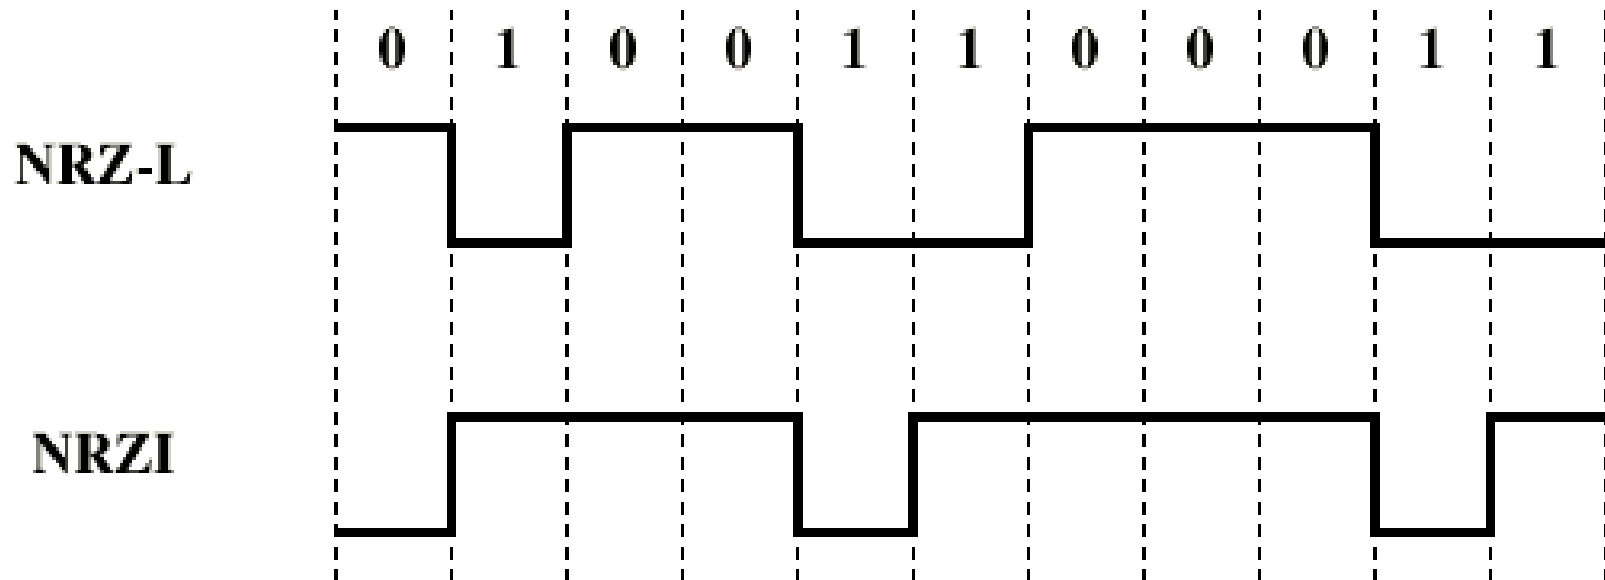
NRZ pros and cons

- Pros
  - Easy to engineer
  - Makes good use of bandwidth
- Cons
  - Presence of DC component
  - Lack of synchronization capability, not self-clocking
    - Long string of 1's (or 0's) for NRZ-L or long string of 0's for NRZI will give constant voltage levels may cause loss of receiver synchronization
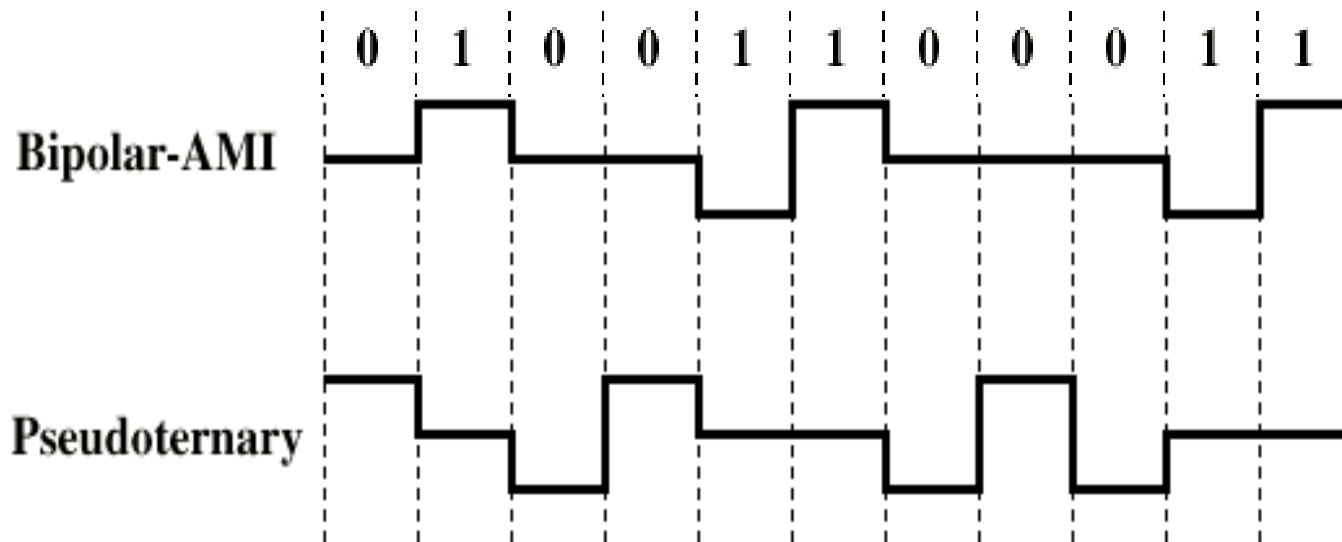
Multilevel Binary

- Use more than two levels
- Bipolar-AMI
  - Bit 0 represented by no line signal
  - Bit 1 represented by positive or negative pulse
  - Pulses for bit 1 alternate in polarity
  - No loss of sync if a long string of 1's happens (long string of 0's still a problem)
  - No net dc component
  - Lower bandwidth
  - Some error detection builtin

Pseudoternary

- Opposite of Bipolar-AMI
- Bit 1 represented by absence of line signal
- Bit 0 represented by alternating positive and negative
- No advantage or disadvantage over bipolar-AMI

Bipolar-AMI and Pseudoternary

Tradeoff for Multilevel Binary

- Not as efficient as NRZ
  - With multi-level binary coding, the line signal may take on one of 3 levels, but each signal element, which could represent $\log_2 3 = 1.58$ bits of information, bears only one bit of information
  - Receiver must distinguish between three levels (+A, -A, 0), increased complexity
 - Requires more signal power for same probability of bit error

Biphase

- Manchester
  - Transition in middle of each bit period
  - Transition serves as clock and data
  - Low to high represents 1
  - High to low represents 0
- Differential Manchester
  - Midbit transition is for clocking only (always there)
  - Transition at start of a bit period represents 0
  - No transition at start of a bit period represents 1
  - Differential encoding scheme

Biphase Pros and Cons

- Con
  - At least one transition per bit time and possibly two
  - Maximum modulation rate is twice that of NRZ
  - Requires more bandwidth
- Pros
  - Synchronization on mid bit transition (self clocking)
  - No dc component
  - Error detection
    - Absence of expected transition points to error in transmission

Digital data, Digital signal

Modulation Rate

- Modulation rate/Signaling rate of an encoding scheme = rate of signal transitions in the channel

- Higher the modulation rate, higher the bandwidth required

- Depends on the actual data stream
  - Should see the maximum modulation rate possible for some data stream

- Another way to characterize : No. of signal transitions per data bit (average/maximum)

5 bits = 5 μs

1   1   1   1   1

NRZI

1 bit =
1 signal element =
1 μs

Manchester

1 bit =
1 μs

1 signal element =
0.5 μs

	Minimum	101010…	Maximum
NRZ-L	0 (all 0s or 1s)	1.0	1.0
NRZI	0 (all 0s)	0.5	1.0 (all 1s)
Bipolar-AMI	0 (all 0s)	1.0	1.0
Pseudoternary	0 (all 1s)	1.0	1.0
Manchester	1.0 (1010…)	1.0	2.0 (all 0s or 1s)
Differential Manchester	1.0 (all 1s)	1.5	2.0 (all 0s)

- So there is a tradeoff
  - Biphase is good for synchronization but needs high signaling rate (higher bandwidth)
  - Non-Biphase codes need lower signaling rate (so lower bandwidth) but may cause receiver to go out of sync
- Can we try to get good clock synchronization at low signaling rate?
  - Scrambling
  - xB/yB codes (for different x, y values)

Scrambling

- Replace sequences that will cause long spells of constant voltage with filling sequences that introduces sufficient transitions

- Receiver should replace filling sequences with original data sequence

- Example: Bipolar-AMI with 8-zeroes substitution (B8ZS)
  - For any octet of all 0's
    - If last voltage pulse before the octet is +ve, replace with 000+–0–+
    - If last voltage pulse before the octet is –ve, replace with 000–+0+–
  - How does the receiver detect scrambling has been done?
    - Hint: Is this a valid sequence for Bipolar-AMI?

8B/10B Encoding

- Other similar schemes like 4B/5B, 5B/6B, 64B/66B (Block Codes)
- Two levels of encoding
  - Bit pattern encoded with another bit pattern
  - Encoded bit pattern encoded again over digital/analog signals (as studied)
- Data split into 8-bit Octets
- Each octet is replaced with a predefined 10-bit code
- Each 10-bit code has one of the following
  - 5 0's and 5 1's
  - 6 0's and 4 1's
  - 4 0's and 6 1's
- Ensures that in the final encoded string, not more than 5 0's or 5 1's can occur simultaneously
  - Ensures enough transitions for clock synchronizations

- Also ensures DC-balance
  - In a string of at least 20 bits, difference in count of 0's and 1's is at most 2
  - This requires introducing *running disparity* bits (we will not cover)
- Requires less b/w, can go over longer distances
- Reduces data rate (as more no. of bits transmitted)
- Used in many protocols, Ex: Gigabit Ethernet

Encoding Digital Data with Analog Signals

- Transmitting digital data through PSTN (Public telephone system)
  - 300Hz to 3400Hz bandwidth
  - modem (modulator-demodulator) is used to convert digital data to analog signal and vice versa
- Carrier signal: basic signal that is changed (modulated)
- Three basic modulation techniques are used:
  - Amplitude shift keying (ASK)
  - Frequency shift keying (FSK)
  - Phase shift keying (PSK)

Amplitude Shift Keying

- Values represented by different amplitudes of a carrier signal
- Usually, one amplitude is zero
  - i.e. presence and absence of carrier is used
- Susceptible to noise
- Used over optical fiber (but usually with some digital encoding before ASK)
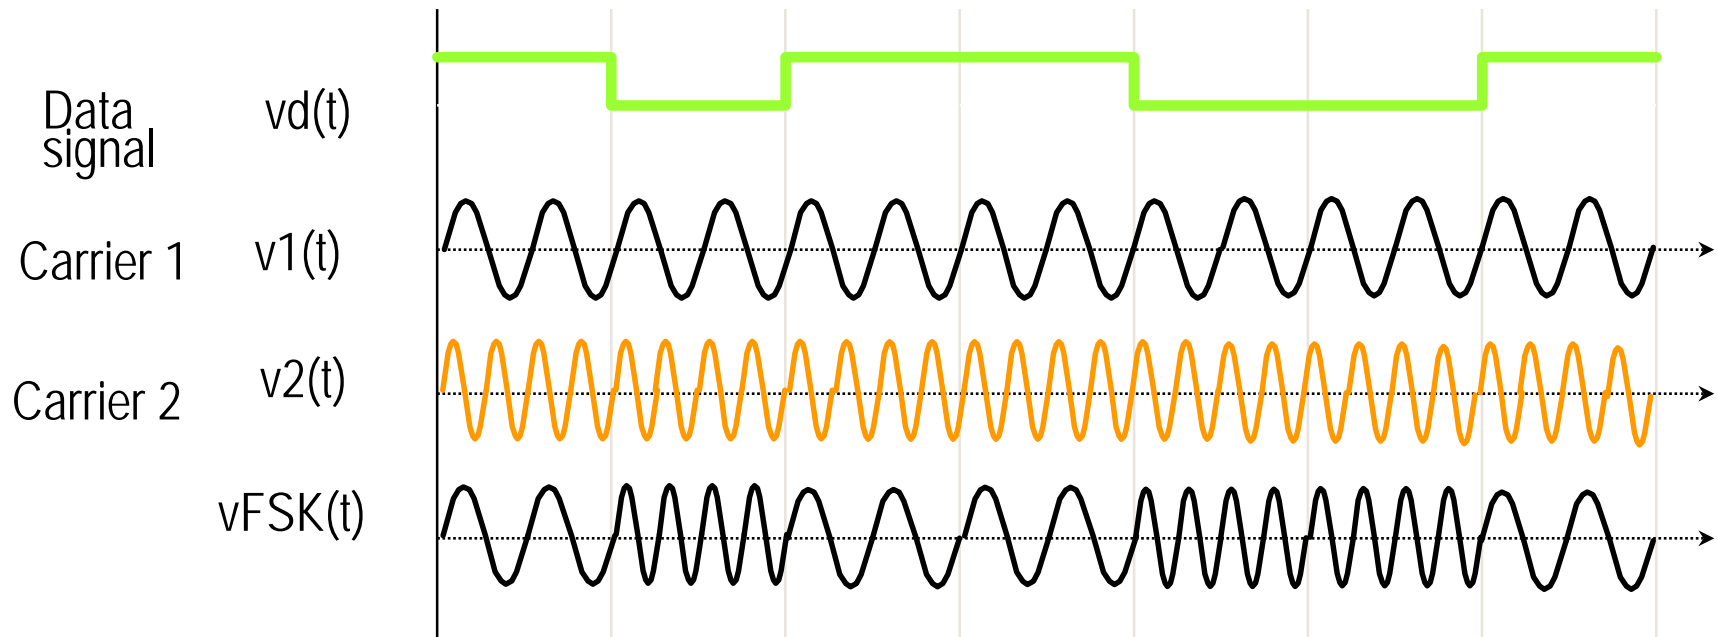
ASK

Vd(t)

Vc(t)

VASK(t)
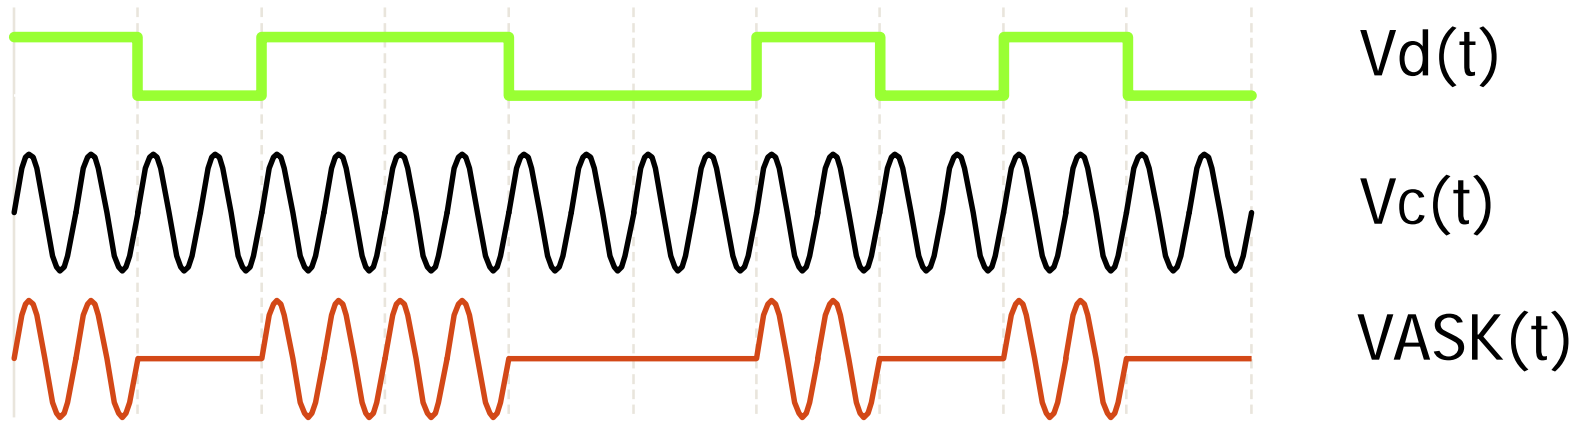
Frequency Shift Keying

- Values represented by different frequencies (near carrier)

- Less susceptible to noise than ASK

- Up to 1200bps on voice grade lines, high frequency radio (3-30 MHz)

- Constrained by bandwidth of channel

FSK

Data signal   vd(t)
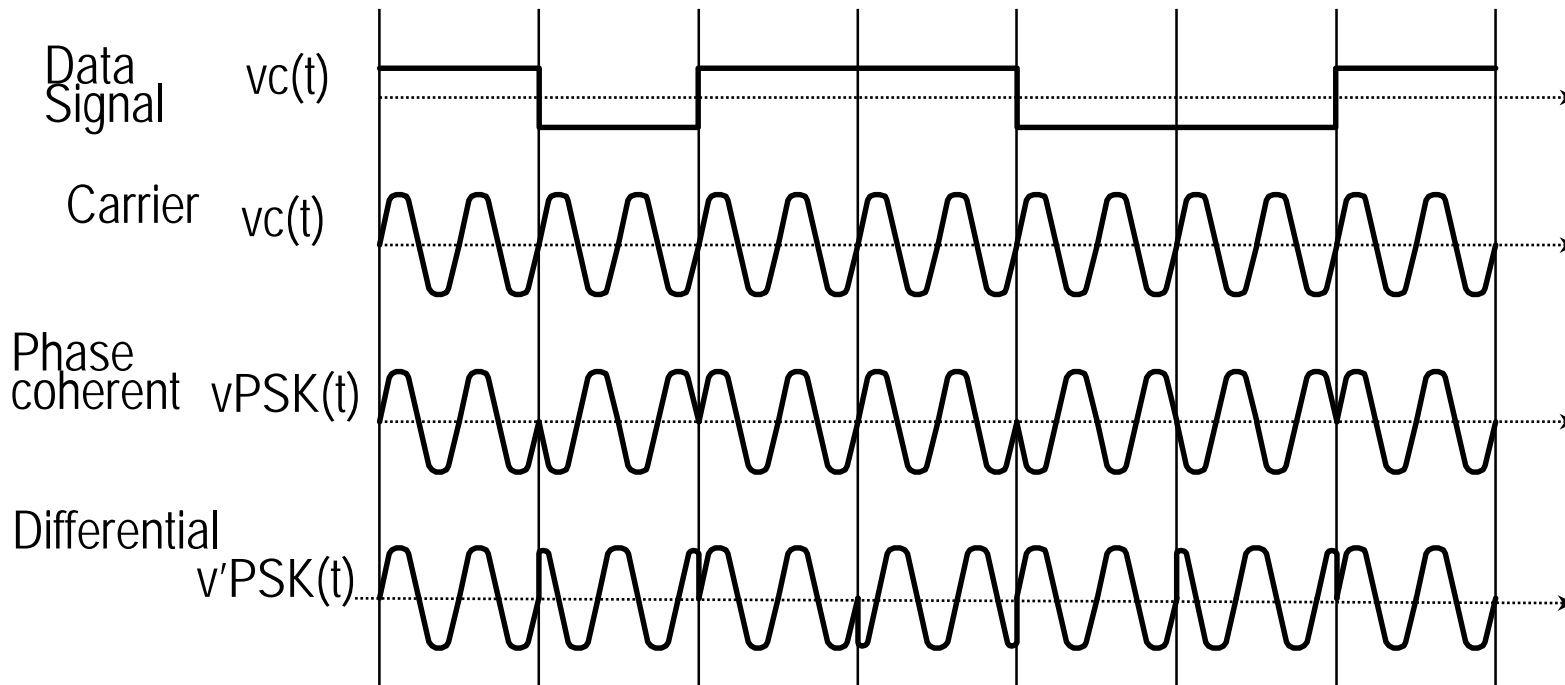
Carrier 1   v1(t)

Carrier 2   v2(t)

vFSK(t)

Phase Shift Keying

- Phase of carrier signal is shifted to represent data

- Differential PSK

  - Phase shifted relative to previous transmission rather than some reference signal

PSK and DPSK

Data Signal vc(t)

Carrier vc(t)

Phase coherent vPSK(t)

Differential v'PSK(t)

Differential example: for every logic 1,
180 degree phase shift

Modulation Techniques

0 0 1 1 0 1 0 0 0 1 0

(a) Amplitude-shift keying

(b) Frequency-shift keying

(c) Phase-shift keying

Quadrature PSK

- More efficient use by each signal element representing more than one bit
    - e.g. shifts of $\pi/2$ (90°)
    - Each element represents two bits
    - Bit rate = 2 x baud rate
- High-speed modems, cellular base stations use different combination of phase and amplitude (*Quadrature Amplitude Modulation or QAM*)
    - 16-QAM (4 amplitude levels and 4 phase levels), allows one (amplitude, phase) combination to encode 4 bits
    - 64-QAM, 256-QAM widely used

Error Detection

Error Detection & Correction

- Error – bit pattern sent and bit pattern received are different. Possible causes
  - Framing errors
  - Noise…
- Error detection – detects error but may or may not correct
- Error correction – detects and corrects errors
  - Usually requires larger number of extra bits with data bits, not usually done in network protocols, we will skip

Error Detection Techniques

- Basic Principle
  - Transmitter: For a given bit stream M, additional bits (called error-detecting bits or check bits) are calculated as a function of M and appended to the end of the data bits in M
  - Receiver: On receiving the bit stream, separates the data bits, performs the same calculation on data bits, and compares the two results. A detected error occurs if there is a mismatch.
- Common Methods
  - Parity Check
  - Cyclic Redundancy Codes
  - Checksum

Parity Check

- One extra "parity" bit is added to each word
  - Odd parity: bit added so as to make # of 1's odd
  - Even parity: makes total # of 1's even
- Detects any odd number of bit errors (error in 1bit, 3 bits, 5 bits …), but can be fooled by any even number of errors
- Simple and easy to implement
- Not very robust against noise

CRC

- Powerful error detection method, easily implemented in hardware

- Message (M) to be transmitted is appended with extra frame checksum bits (F), so that bit pattern transmitted (T) is perfectly divisible by a special "generator" pattern (P) - (divisor)

- At destination, divide received message by the same P. If remainder is nonzero, there is an error

- Let
  - T = (k+n)-bit frame to be transmitted, n<k
  - M = k-bit message, the first k bits of T
  - F = n-bit FCS, the last n bits of T
  - P = n+1 bits, generator pattern (predetermined divisor)
- The concept uses modulo-2 arithmetic
  - no carries/borrows; add ≡ subtract ≡ XOR

- Extend M with n '0's to the right ($\equiv 2^n M$)(shift left by n bits)
- Divide extended message by P to get R ($2^n M / P = Q + R/P$)
- Add R to extended message to form T ($T = 2^n M + R$)
- Transmit T
- At receiver, divide T by P. Nonzero remainder $\Rightarrow$ error

$$\frac{T}{P} = \frac{2^n M + R}{P} = Q + \frac{R}{P} + \frac{R}{P} = Q + \frac{R+R}{P} = Q$$

Note: R+R=0 in mod-2 arithmetic

$$0 \oplus 0 = 0$$
$$1 \oplus 1 = 0$$

Note: Remainder R=F=FCS in these examples

- M = 110011, P = 11001, R = 4 bits
- Append 4 zeros to M, we get 1100110000

```
– 2.
                              1 0 0 0 0 1
         1 1 0 0 1 ) 1 1 0 0 1 1 0 0 0 0
                   ⊕ 1 1 0 0 1
                     0 0 0 0 1
                   ⊕ 0 0 0 0 0
                     0 0 0 1 0
                   ⊕ 0 0 0 0 0
                     0 0 1 0 0
                   ⊕ 0 0 0 0 0
                     0 1 0 0 0
                   ⊕ 0 0 0 0 0
                     1 0 0 0 0
                   ⊕ 1 1 0 0 1
                     1 0 0 1
```

For each stage of division, if the number of dividend bits equals number of divisor P bits, then Q=1, otherwise Q=0

– 3. ∴ T = 110011 1001

- Can view CRC generation in terms of polynomial arithmetic also

  - Any bit pattern ≡ polynomial in dummy variable X with the bits as coefficients as shown in the following example:

    - e.g., M = 110011 ≡ $1 \cdot X^5 + 1 \cdot X^4 + 0 \cdot X^3 + 0 \cdot X^2 + 1 \cdot X + 1 \cdot X^0$
      $\therefore$ M(X) = $X^5 + X^4 + X + 1$

- CRC generation in terms of polynomial
  - Append n '0's ≡ $X^n M(X)$
  - Modulo 2 division → $\dfrac{X^n M(X)}{P(X)} = Q(X) + \dfrac{R(X)}{P(X)}$

  - Transmit $X^n M(X) + R(X) = T(X)$
  - At receiver:

$$\frac{X^n M(X) + R(X)}{P(X)} = \frac{X^n M(X)}{P(X)} + \frac{R(X)}{P(X)}$$

$$= Q(X) + \boxed{\frac{R(X) + R(X)}{P(X)}}$$

→ If no error 0 always

- Commonly used polynomials, P(X)
  - CRC-16 = $X^{16} + X^{15} + X^2 + 1$
  - CRC-CCITT = $X^{16} + X^{12} + X^5 + 1$
  - CRC-32 = $X^{32} + X^{26} + X^{23} + X^{22} + X^{16} + X^{12} + X^{11} + X^{10} + X^8 + X^7 + X^{12} + X^4 + X^2 + X + 1$

- CRC can detect
  - All single-bit errors
  - All double-bit errors, as long as $P(X)$ has a factor with at least three terms (as long as P has at least three 1s)
  - Any odd number of errors, as long as $P(X)$ contains a factor $(X+1)$
  - Any burst error for which the length of the burst is less than or equal to the length of the FCS
  - Many other larger burst errors
- Strong error detection scheme, easy implementation in hardware

Checksum

- Another common error detection scheme
- Break up the bit streams into fixed length "words" (typically 16 bit)
- Transmitter
  - Do 1's complement addition of all the words
  - Take the 1's complement of the final sum and send it with the bit stream
- Receiver
  - Do one's complement addition of all the words (including the checksum
  - Should be all 1's

Example

- Input stream (in hex): 00 01 F2 03 F4 F5 F6 F7

- Stream sent : 00 01 F2 03 F4 F5 F6 F7 22 0D

Partial sum	0001 F203 F204		Partial sum	0001 F203 F204
Partial sum	F204 F4F5 1E6F9		Partial sum	F204 F4F5 1E6F9
Carry	E6F9 1 E6FA		Carry	E6F9 1 E6FA
Partial sum	E6FA F6F7 1DDF1		Partial sum	E6FA F6F7 1DDF1
Carry	DDF1 1 DDF2		Carry	DDF1 1 DDF2
Ones complement of the result	220D		Partial sum	DDF2 220D FFFF

(a) Checksum calculation by sender   (b) Checksum verification by receiver

- Better than parity check
  - Higher error detection capability
- Less effective than CRC
- Still used a lot in higher layer protocols because of simplicity and low overhead of computing the checksum

Flow Control

Flow Control

- What if sender sends the bit pattern too fast? Receiver may not be able to process the data as fast as the sender is sending it
  - Receiver can buffer, but buffer has finite size. After that is filled, data is lost
- Flow control – technique to control data flow between sender and receiver so that sender is blocked if receiver cannot accept any more data
  - Sender must get an acknowledgement from the receiver before it can send more data

Stop & Wait Flow Control

- Sender sends a frame

- Receiver receives frame & acknowledges it

- Sender waits to receive "ack" before sending next frame (If receiver is not ready to receive another frame it holds back the ack)

- One frame at a time is sent over the transmission line

- Simple, easy to implement

- Problems:
  - Works fine if there is only a few large frames to be sent
  - Usually we want smaller frames
    - Receiver's buffer size may be small
    - If a large frame is transmitted, the entire frame has to be retransmitted if there is an error
      - Longer the frame, larger the chance of an error
    - On a shared medium, do not want any one sender to occupy the medium for a long time (will see later)
  - With smaller frames, takes too long to send all frames

Link Efficiency of Stop & Wait

- Bit Length of link in bits B = maximum possible number of bits that can be simultaneously present in the link
  - $B = R \times (D/V)$
    - R = data rate of link in bps
    - D = length of link in meters
    - V = propagation speed in m/sec
- So for Stop & Wait, ideally we want frame length to be close to bit length of the link
  - The sender has enough bits to put out on the link while the earlier bits of the same frame reach the receiver
  - Links stays occupied for more time
- But we have also argued that we want frame sizes to be small
- Results in very poor line/link utilization for Stop & Wait, especially in fast links over longer distances
  - B becomes very large, so impractical to have such large frames

- Consider a frame length of L bits
- Transmission delay = Time to put the frame into the link = L/R
- Time for last bit of the frame to reach the receiver = (L/R) + (D/V)
- Assume that processing delay at receiver is negligible
- Assume that transmission time of ACK is negligible (very small frame)
- Time for ACK to come to sender = (D/V)
- Total time to transmit the frame T = (L/R) + 2(D/V)
- Total no. of bits that could have ben transmitted in this time = T × R
- Actual no. of bits transmitted = L
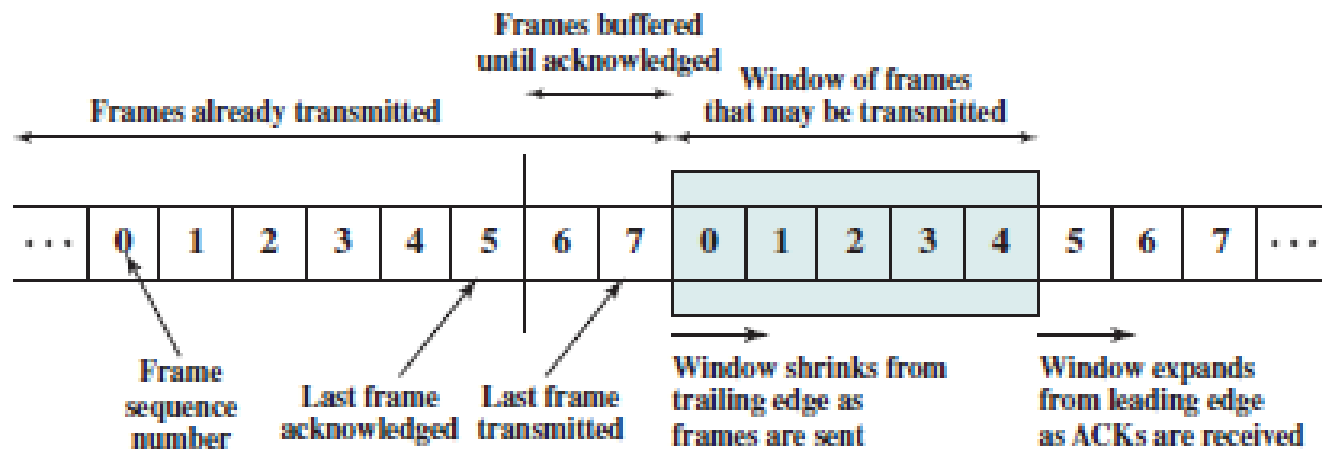- Line utilization = L/(T × R) = L / (L + 2R(D/V))

Sliding Window Flow Control

- Reduces line inefficiency problem of stop-and-wait by transmitting multiple frames without waiting for acknowledgement

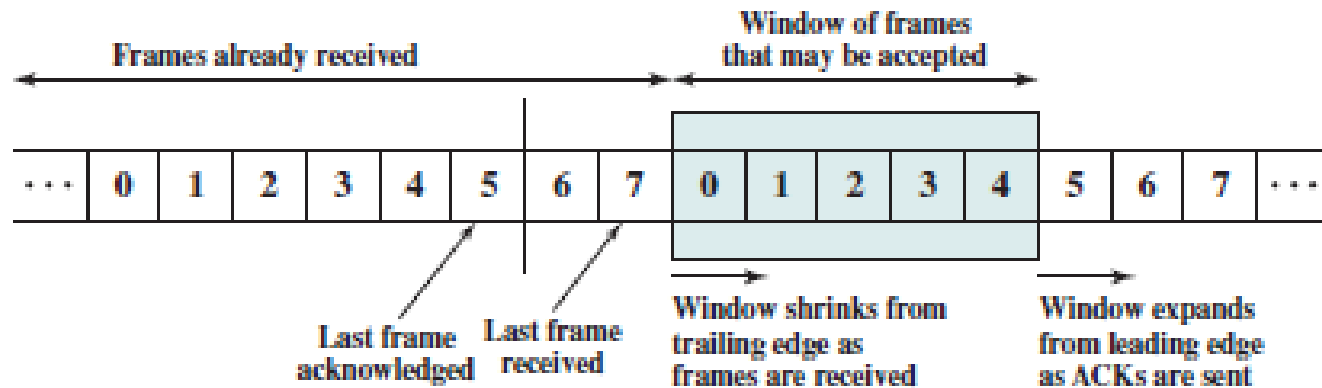- Suppose two stations A and B are connected by a full-duplex link

A ←——————————————————————→ B

Sliding Window (contd.)

- Station B allocates buffer space for W frames.
  - Thus B can accept W frames, and A is allowed to send W frames without waiting for any acknowledgement
- Sender maintains a list of sequence numbers that it is allowed to send (sender window)
- Receiver also maintains a list of sequence numbers that it is prepared to receive (receiver window)
- Since the sequence number to be used occupies a field in the frame, it is clearly of bounded size
  - For a k-bit field, the range of sequence numbers is 0 through 2^k-1, and frames are numbered modulo 2^k
  - Maximum window size $= 2^k - 1$ (why?)

- Each frame is labeled with a sequence number
  - To keep track of the frames which have been acknowledged
  - B acknowledges a frame by sending an ACK/RR that includes the sequence number of the next frame expected. This also explicitly announces that B is prepared to receive the next W frames, beginning with the number specified
- This scheme can be used to acknowledge multiple frames
  - B could receive frames 2,3,4 but withhold ACK until frame 4 has arrived. By returning an ACK with sequence number 5, B acknowledges frames 2,3,4 at one time

Frames buffered until acknowledged

Window of frames that may be transmitted

Frames already transmitted

... | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | ...

Frame sequence number

Last frame acknowledged   Last frame transmitted

Window shrinks from trailing edge as frames are sent

Window expands from leading edge as ACKs are received

(a) Sender's perspective

Frames already received

Window of frames that may be accepted

... | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | ...

Last frame acknowledged   Last frame received

Window shrinks from trailing edge as frames are received

Window expands from leading edge as ACKs are sent

(b) Receiver's perspective

Sliding Window (contd.)

- The actual window size need not be the maximum possible size for a given sequence number length
  - For a 3-bit sequence number, a window size of 5 can also be configured
- If two stations exchange data, each need to maintain two windows. To save communication capacity, a technique called piggybacking is used
  - Each data frame includes a field that holds the sequence number of that frame plus a field that holds the sequence number used for ACK
  - If a station has an ACK but no data to send, it sends a separate ACK frame

Error Control

Error Control

- Ensures finally received bit pattern is same as sent bit pattern, though may require more than one transmission
- Type of frame loss
  - Lost frames (How can it be lost completely??)
  - Damaged Frames (error detected)
- Two types of error control
  - Forward error control:
    - Error recovery by correction at the receiver [Forward Error Correction (FEC)]
    - Requires extra error correcting bits to be added to data bits (Ex. Hamming code)
    - No. of bits needed large for even small number of bit errors
  - Backward error control:
   - Error recovery by retransmission [Automatic Repeat Request (ARQ)]

ARQ

- Basic idea
  - Break up bit pattern into fixed length frames
  - Add error detection bits with each frame
  - Ask transmitter to retransmit frames that show error or that are expected but not received within a timeout
  - Reconstruct final received bit pattern when all frames received successfully

Stop & Wait ARQ

- Sender transmits message frame

- Receiver checks received frame for errors, sends ACK

- Sender waits for ACK to send next frame

- Two possibilities:
  - Receiver does not receive the frame or receives a damaged frame
    - Receiver simply discards a damaged frame
    - Sender will timeout, resend the frame
    - This continues until an ACK is received

Stop & Wait ARQ (contd.)

- Receiver receives the frame but the ACK is lost or damaged
- Sender times out, or discards the ACK, sends the frame again
- So receiver may receive the frame more than once

- Handling duplicate frames
  - Add a 1-bit frame number, 0 or 1, to each frame
  - ACKS includes a frame number – ACK0 and ACK1
    - ACK0 implies sender has received frame 1 successfully and vice-versa
  - Only need a 1-bit frame number alternating 1 and 0 since they are sent one at a time
    - Two frames are never together on the line

- Pros
  - Simple to implement
  - Only 1-bit frame number needed
- Con
  - Extremely inefficient (as already seen in flow control)
    - Only one frame can be sent at a time
    - Sender has to wait for an ack even if it has more frames to send and even if the frame is received successfully

Go-back-N ARQ

- The sender can send a sequence of frames without waiting for others

- The frames are numbered sequentially

- Window size – the maximum number of unacknowledged frames that can be there

- Bahavior on no error – similar to sliding window

- If the receiver detects an error on a frame, it sends a NACK/REJ for that frame. The receiver will discard all future frames received until the frame in error is correctly received

- The sender, when it receives a NACK/REJ or times out, must retransmit the frame in error plus all succeeding frames. (Sender must maintain a copy of each unacknowledged frame)

- Maximum window size with k-bit sequence number is $2^k - 1$

- What if not?
  - Suppose window size is 8 (= 2^k) for k = 3
  - Sender sends frame 0, receives ACK/RR 1
  - Sender sends frames 1, 2, 3, 4, 5, 6, 7, 0 (next 8 frames)
  - Sender gets another ACK/RR1
  - Ambiguity:
    - Did all 8 frames get through successful?
    - Did all 8 frames were lost/damaged?

Selective-Reject ARQ

- The only frames retransmitted are those that receive a NACK or which timeout

- Pro: Can save retransmissions by sending only frames that are in error

- Con:
  - Requires more buffer space, as receiver must keep enough buffer space to save frames received successfully after the frame in error
  - More complicated logic to reinsert the lost frame in the sequence when received successfully later
  - Lower window size of 2^{k-1} (why?)

Summary

- We now know how to make only two machines communicate connected directly with a link
  - Encode data to be send with signals
  - Break data into frames
  - Make sure sender and receiver are synchronized for proper frame receive
  - Detect errors in frames (error detection) as well as lost frames (timeout)
  - Use error control techniques to ensure all frames (and therefore all data) are received at the receiver correctly eventually
  - Ensure fast sender cannot overwhelm a slow receiver
- Next step: More than one machine sharing a direct link